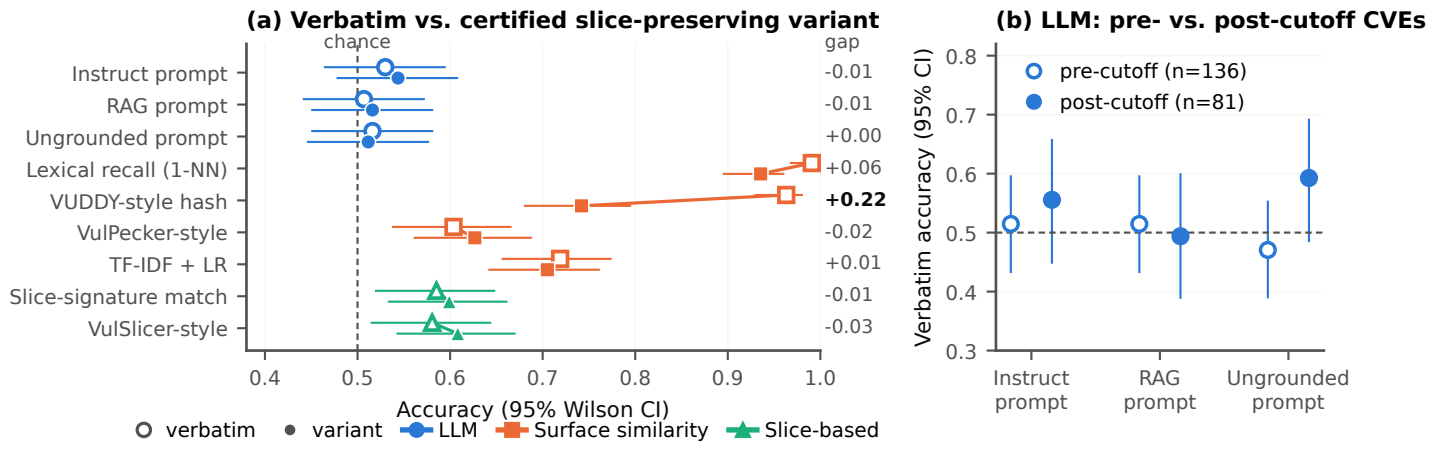
(a) Verbatim vs. certified slice-preserving variant
chance
gap
Instruct prompt
RAG prompt
Ungrounded prompt
Lexical recall (1-NN)
VUDDY-style hash
VulPecker-style
TF-IDF + LR
Slice-signature match
VulSlicer-style
-0.01
-0.01
+0.00
+0.06
+0.22
-0.02
+0.01
-0.01
-0.03
0.4
0.5
0.6
0.7
0.8
0.9
1.0
Accuracy (95% Wilson CI)
verbatim
variant
LLM
Surface similarity
Slice-based
(b) LLM: pre- vs. post-cutoff CVEs
pre-cutoff (n=136)
post-cutoff (n=81)
Verbatim accuracy (95% CI)
0.8
0.7
0.6
0.5
0.4
0.3
Instruct prompt
RAG prompt
Ungrounded prompt